


Project Phoenix

App modernization using K8s on Azure
github.com/CSA-OCP-GER/phoenix



Implementation challenges



Cost



Scale



Complexity

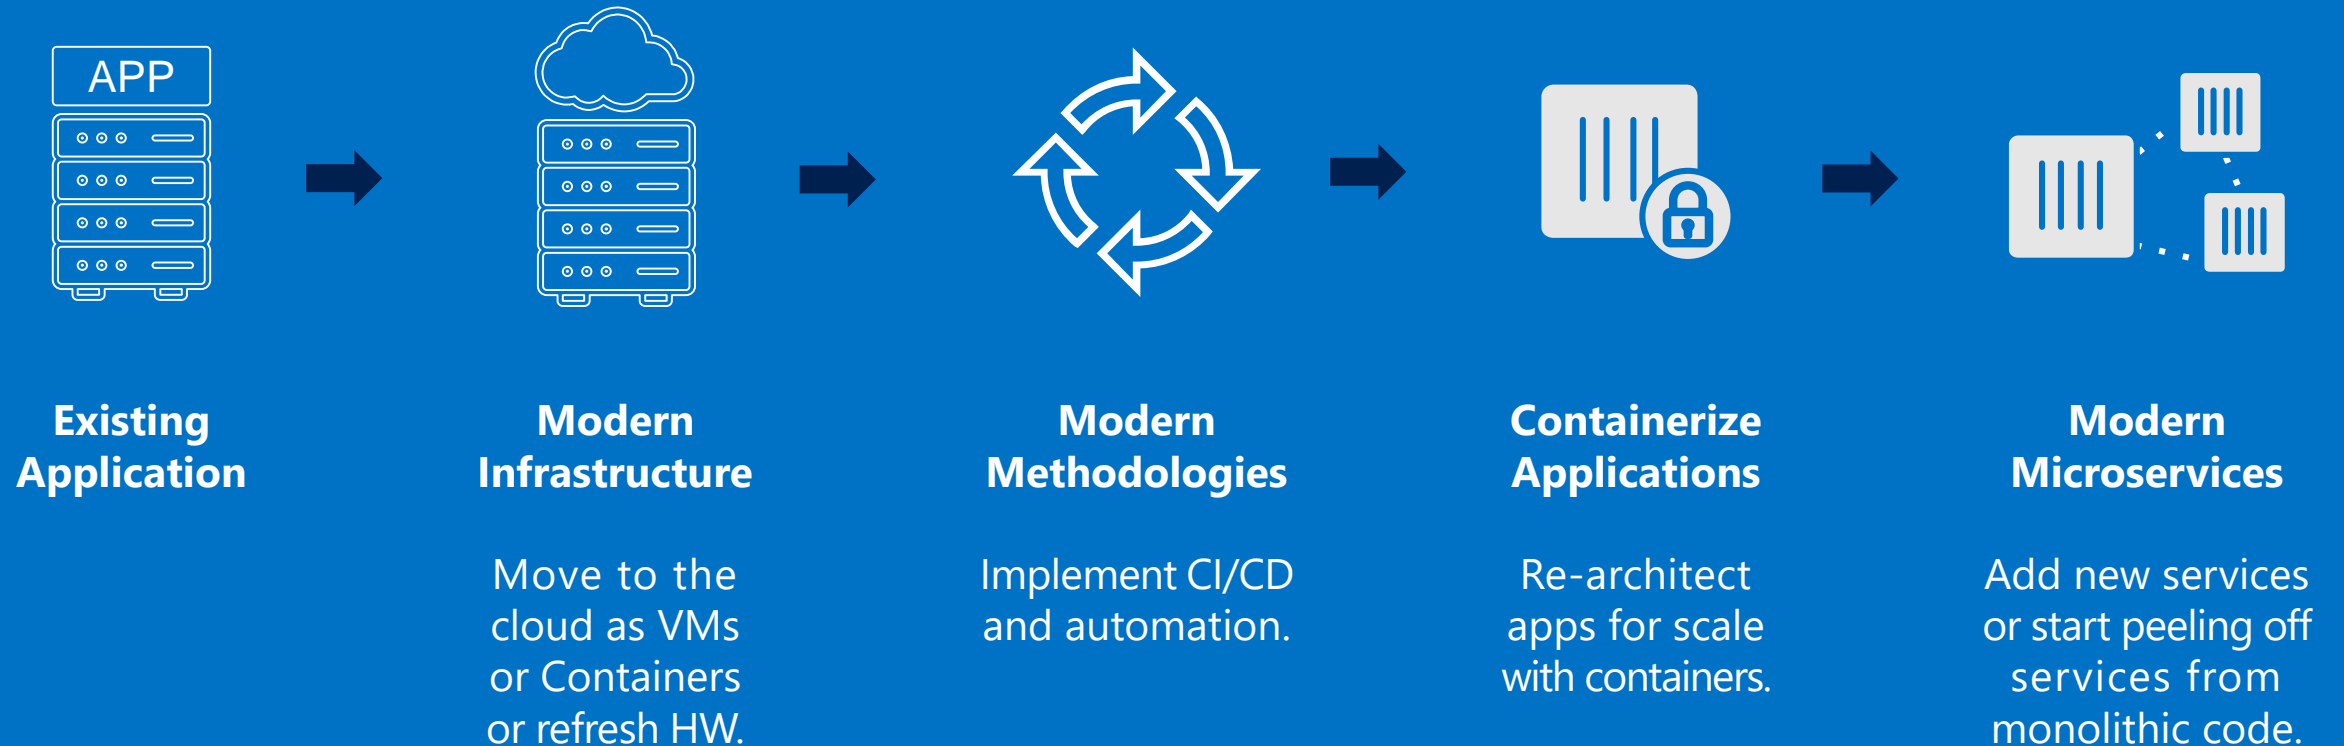


Agility




Security

From traditional app to modern app




Microservices as an architecture pattern

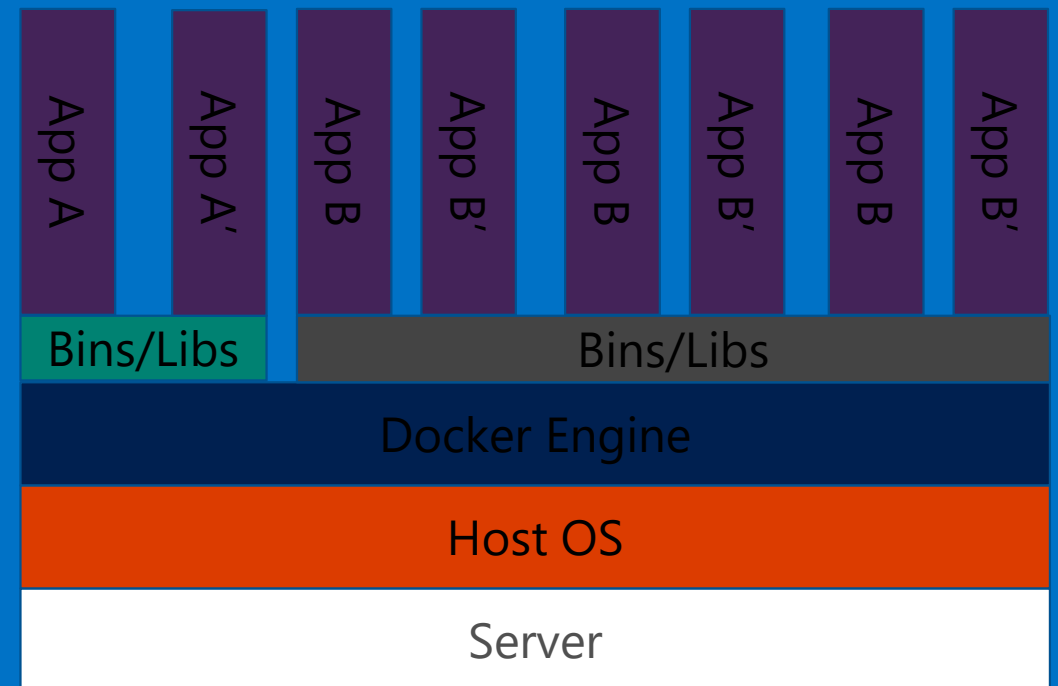
- Individually built and deployed
- Small, independently services
- Integrate using published API
- Fine-grained, loosely coupled



Container 101



Containers are isolated, but share OS and, where appropriate, bins/libraries



microservices \neq containers

microservices is an architectural
design approach

containers are an implementation
detail that often helps

Scaling microservices is hard

450 microservices




500+ microservices



NETFLIX

500+ microservices



Source:

Netflix: <http://www.slideshare.net/BruceWong3/the-case-for-chaos>

Twitter: <https://twitter.com/adrianco/status/441883572618948608>

Hail-o: <https://sudo.hailoapp.com/services/2015/03/09/journey-into-a-microservice-world-part-3/>

Orchestrators = containers at scale

Cluster Management Deploy and manage cluster resources	Scheduling When containers run	Lifecycle & Health Keep containers running despite failure	Naming & Discovery Where are my containers	Load Balancing Distribute traffic evenly
Scaling Make container sets elastic in number	Image Repository Centralized, secure container images	Continuous Delivery CI/CD pipeline and DevOps workflow	Logging & Monitoring Track events in containers and cluster	Storage Volumes Persistent data for containers

12-Factor Apps (1-5)

1. Single root repo; don't share code with another app
2. Deploy dependent libs with app
3. No config in code; read from environment vars
4. Handle unresponsive app dependencies robustly
5. Strictly separate build, release, & run steps
 - Build: Builds a version of the code repo & gathers dependencies
 - Release: Combines build with config ReleaseId (immutable)
 - Run: Runs app in execution environment

12-Factor Apps (6-12)

- 6. App executes as 1+ stateless process & shares nothing
- 7. App listens on ports; avoid using (web) host
- 8. Use processes for isolation; multiple for concurrency
- 9. Processes can crash/be killed quickly & start fast
- 10. Keep dev, staging, & prod environments similar
- 11. Log to stdout (dev=console; prod=file & archived)
- 12. Deploy & run admin tasks (scripts) as processes

Azure Container support



Azure Container Instance



Container Service



Service Fabric

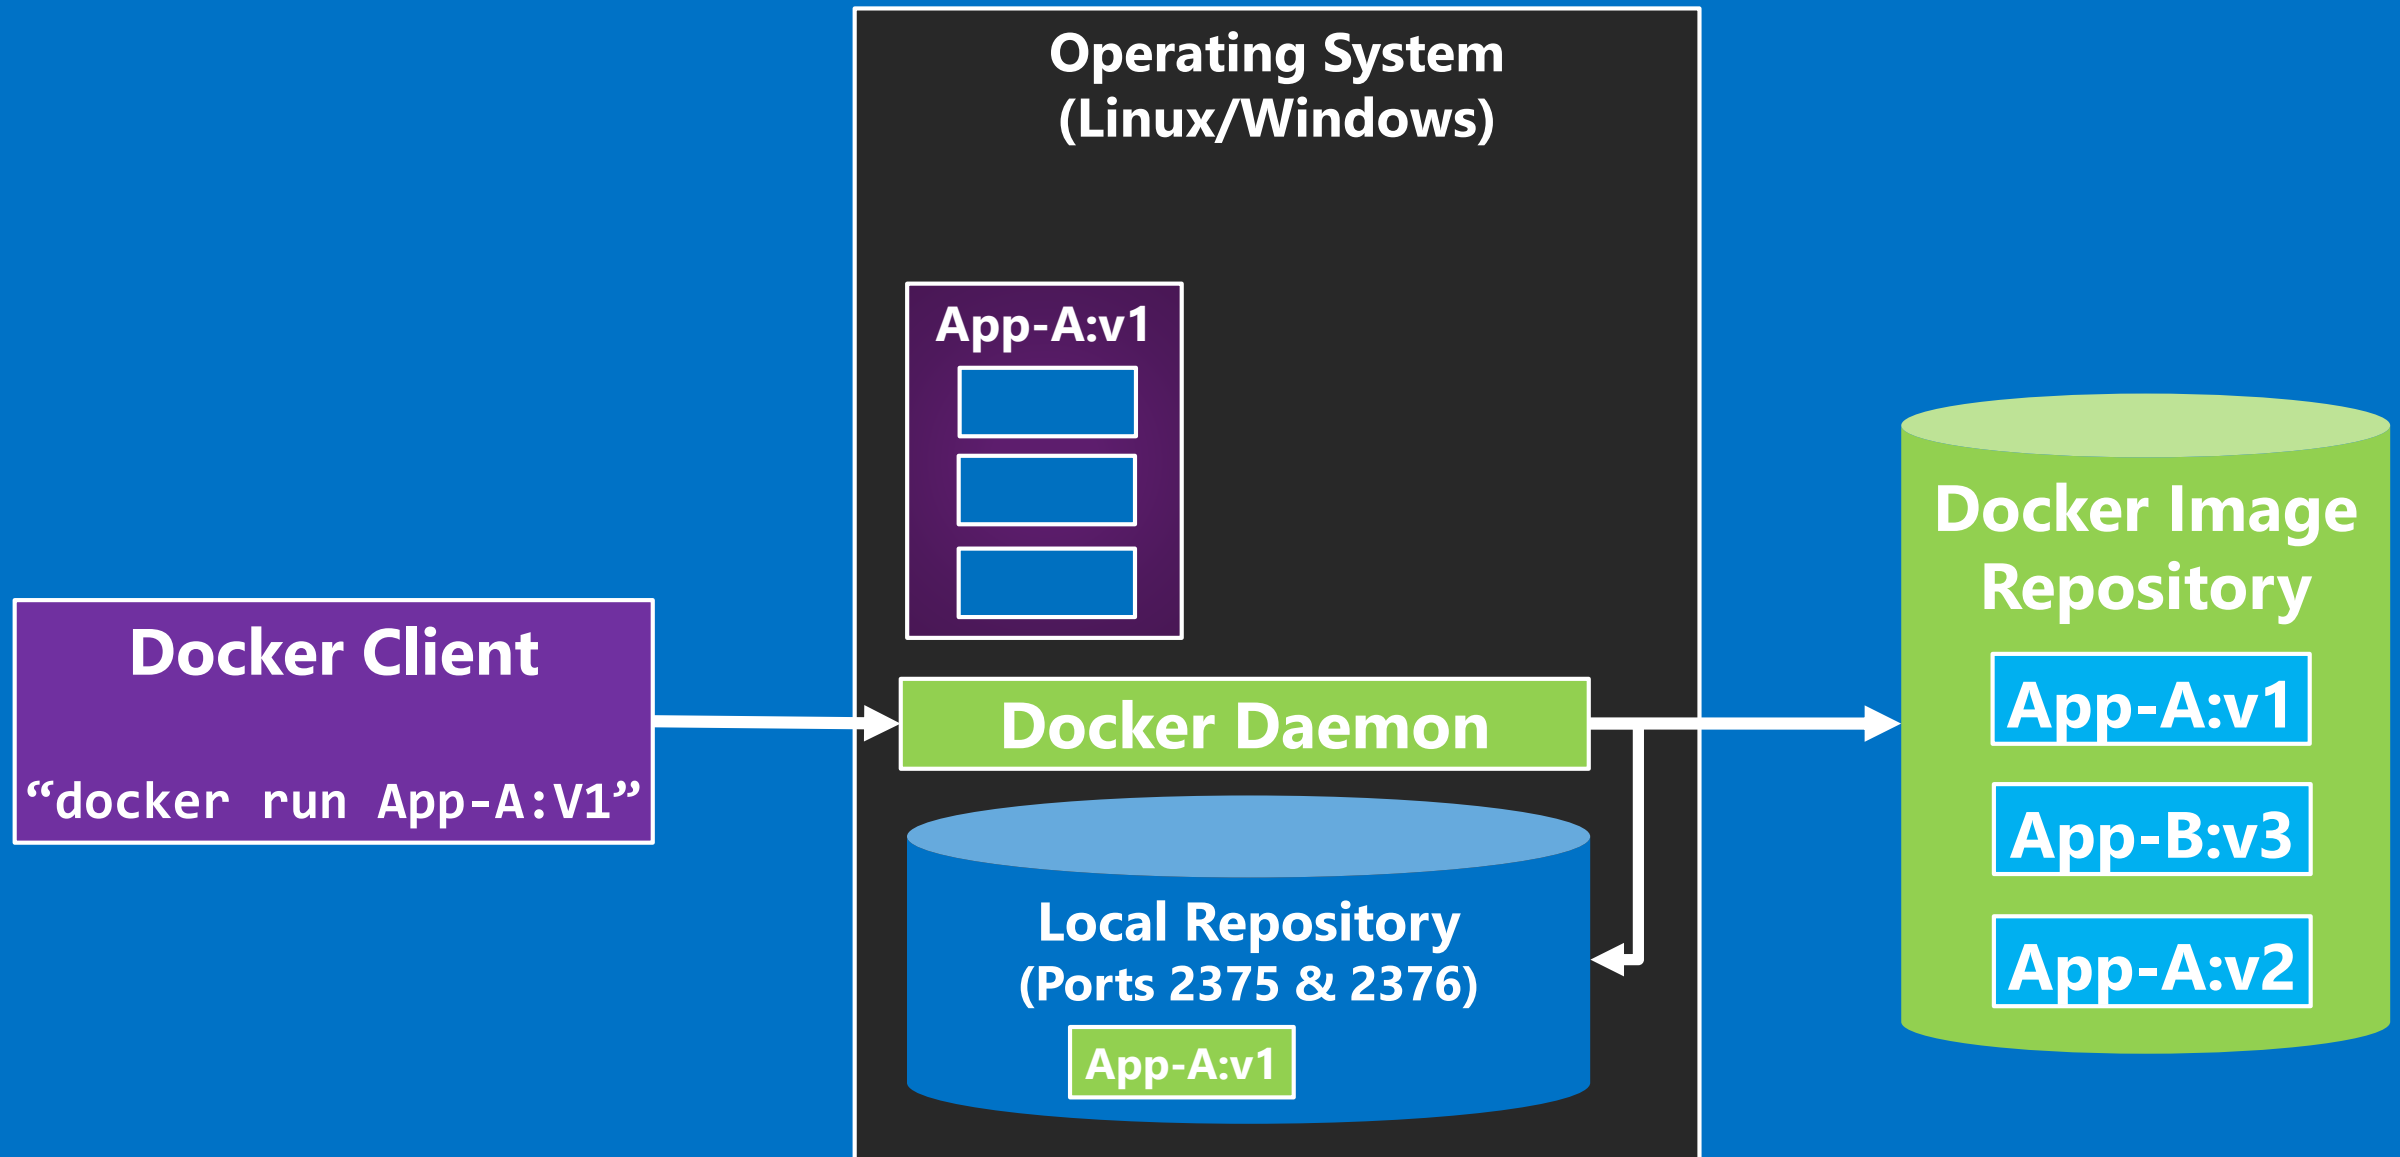


Web Apps

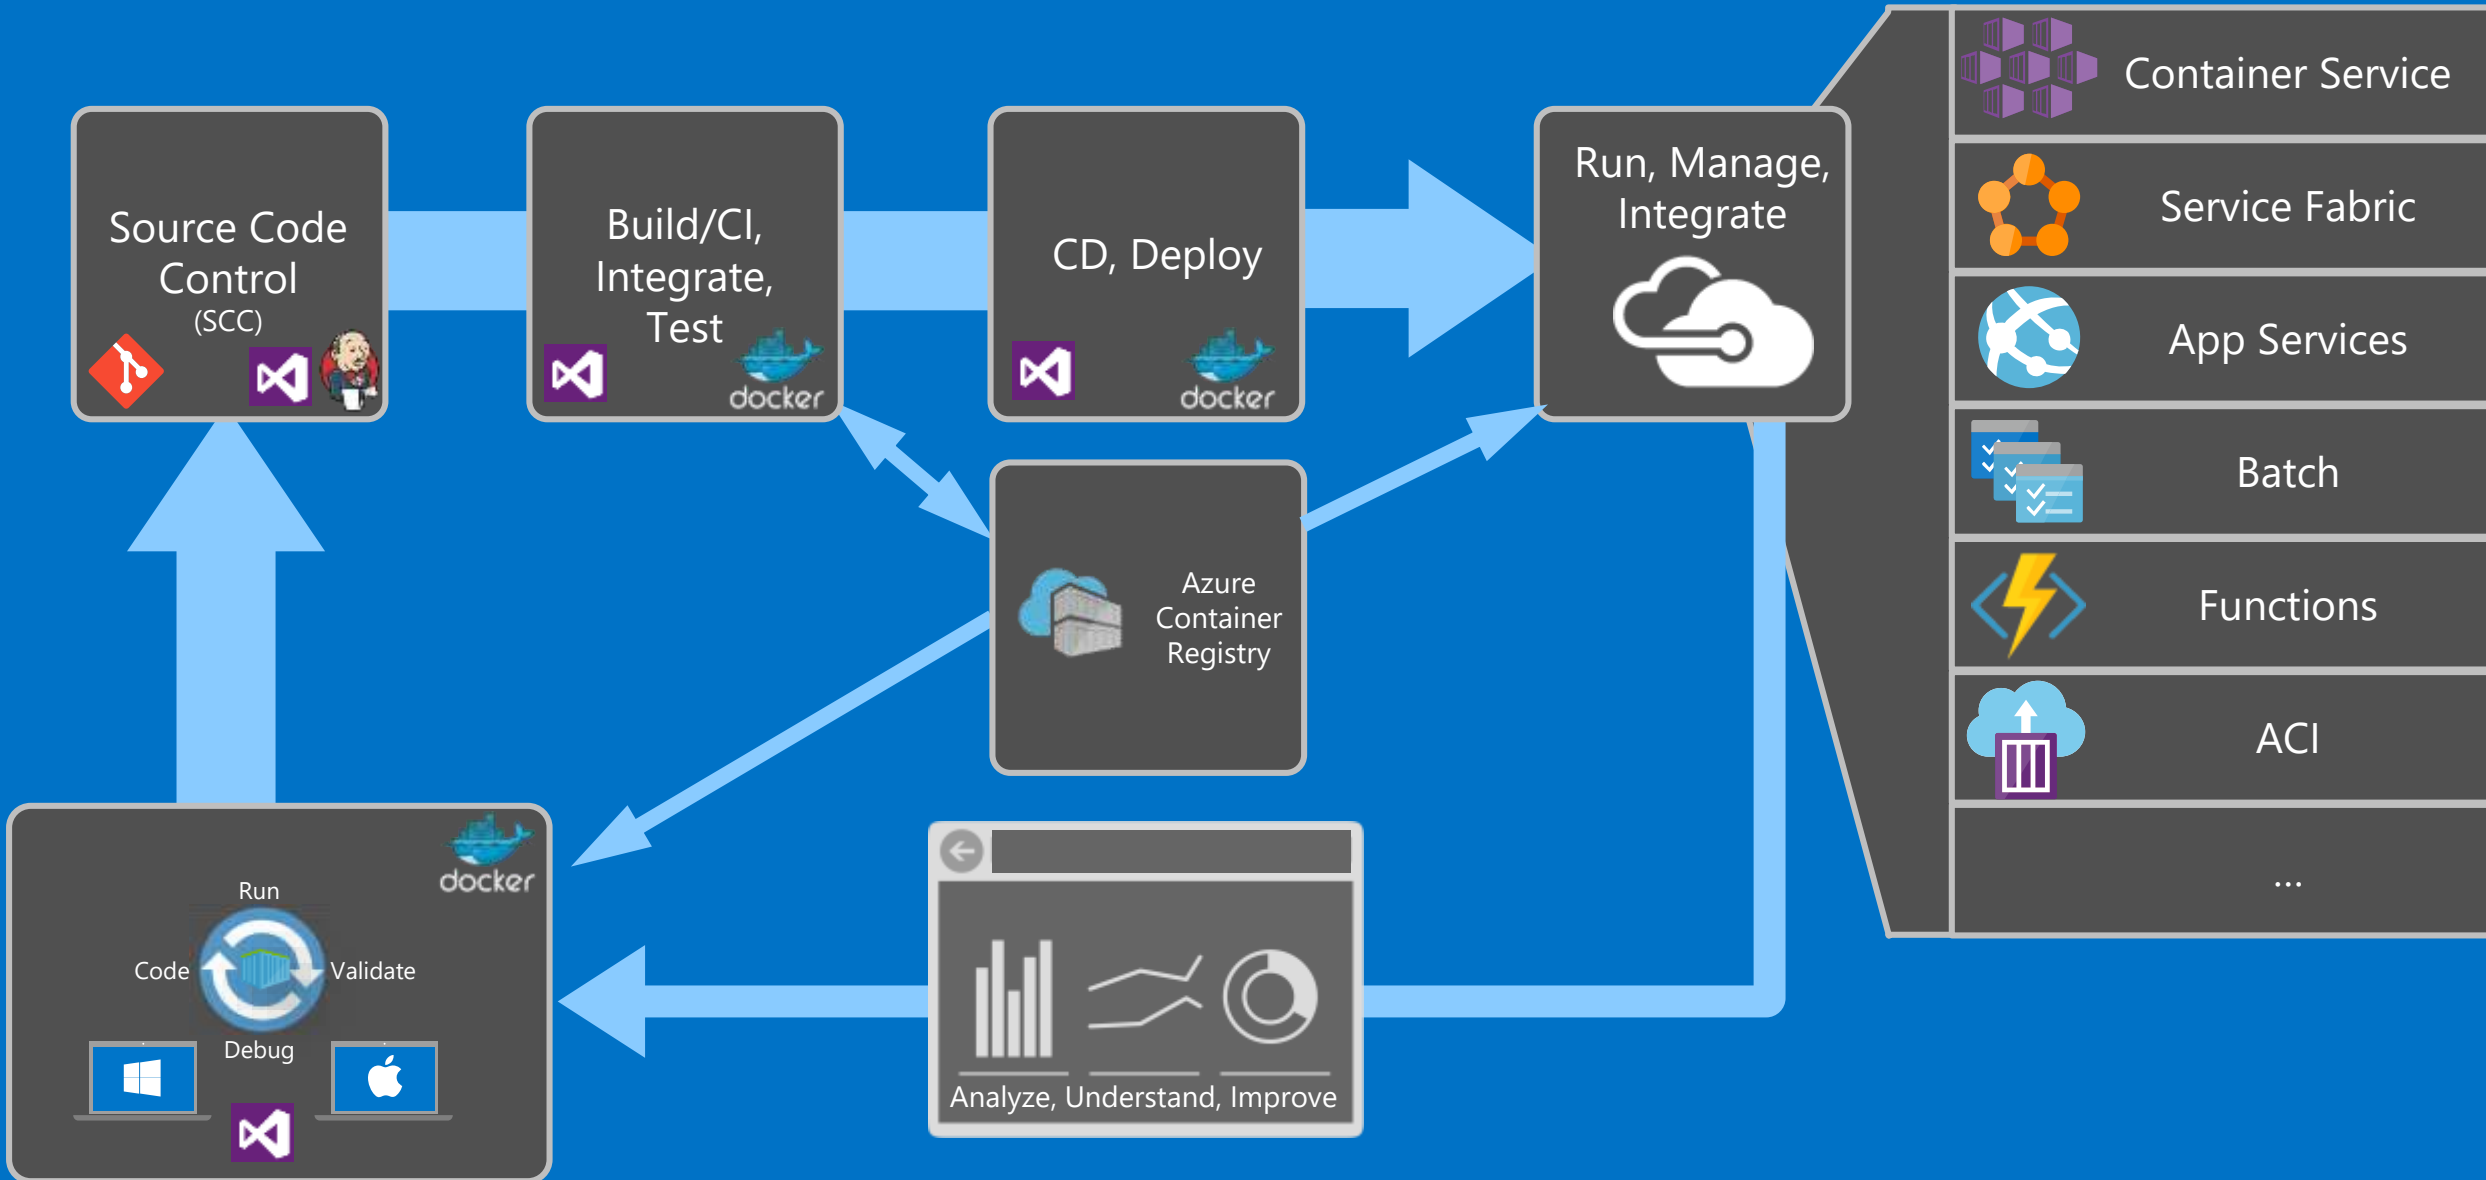


Batch

~\$ docker run App-A:v1



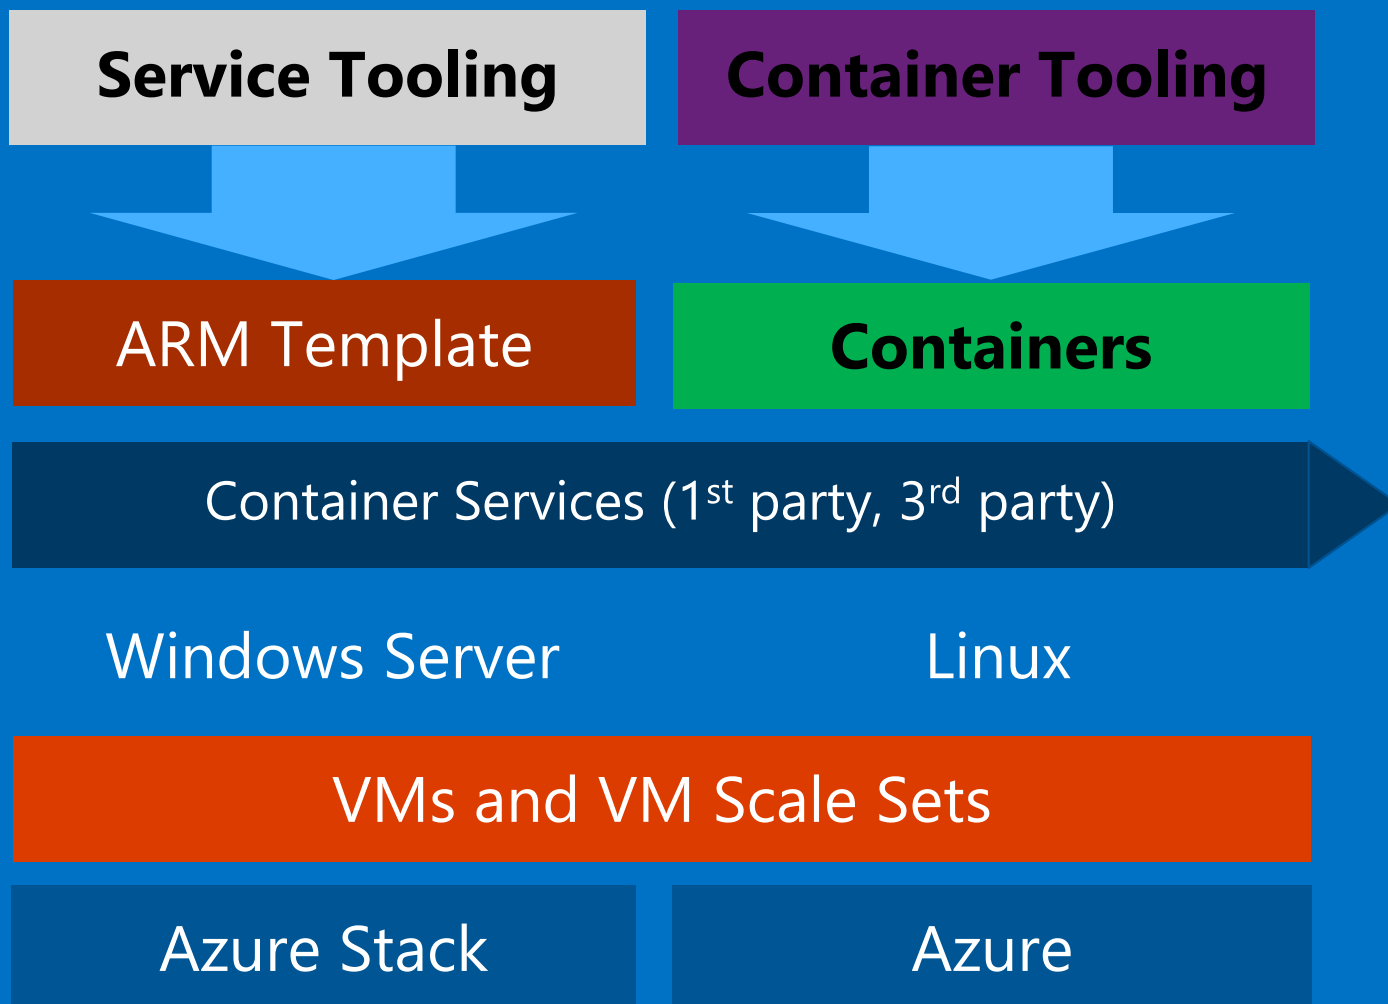
DevOps Pipeline



Azure Container Service



Azure Container Services




Layer	Supported Technologies
Configuration as Code	ARM, Dockerfile, Docker Compose, Marathon.json
Host cluster management	VM Scale Sets
Container orchestration	Docker Swarm, DC/ OS, Kubernetes
Monitoring	OMS, App Insights

Azure Container Service

- ✓ Container hosting solution optimized for Azure
- ✓ The cloud's most open option for containers
- ✓ Manage container applications using familiar tools

Azure



DC/OS



Swarm



Kubernetes

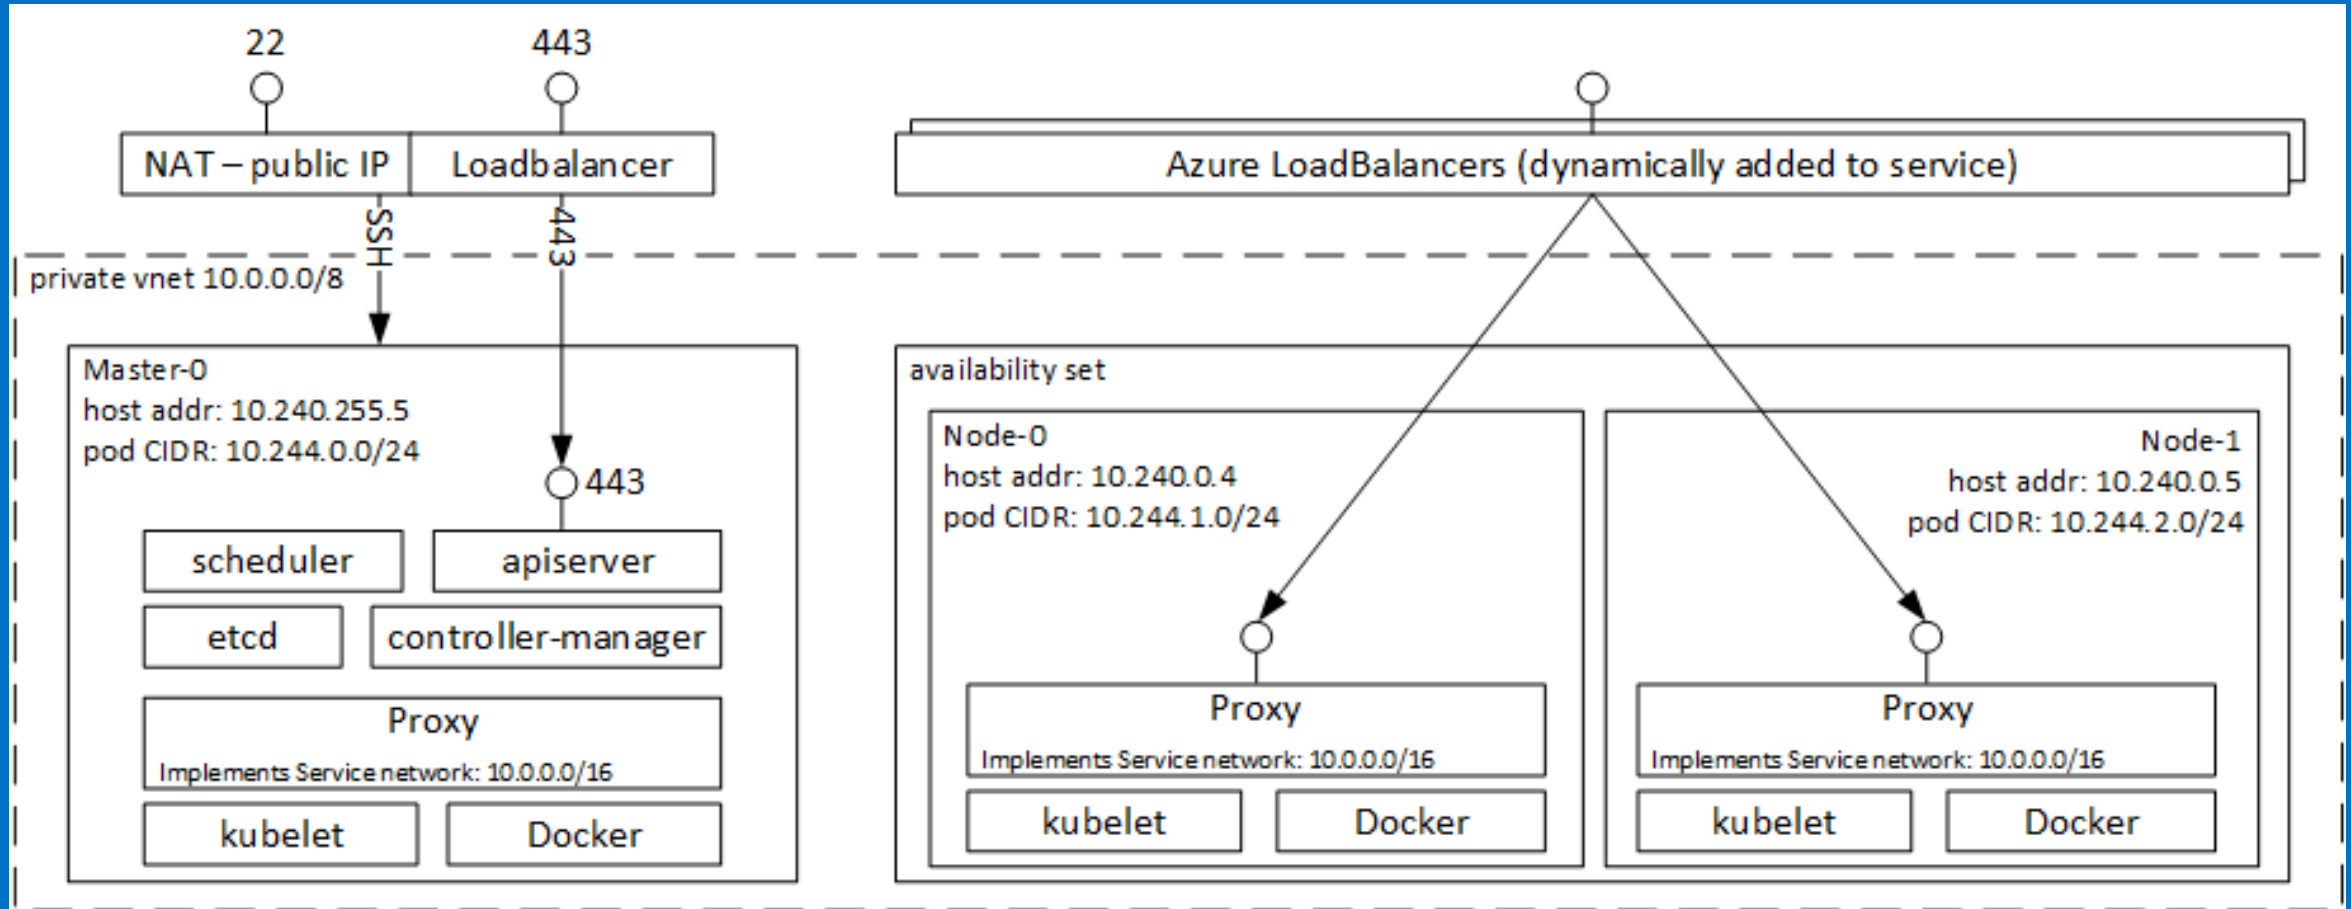
Application design



Application Oriented API



Kubernetes infrastructure setup on Azure



Kubernetes Objects

- Pods
 - Services
 - Deployments
 - Replica Sets
 - ...
- Defined via *.yaml *.yml file
 - Desired state is specified via „Spec“ section
 - Actual state is represented in „Status“ section

Kubelet

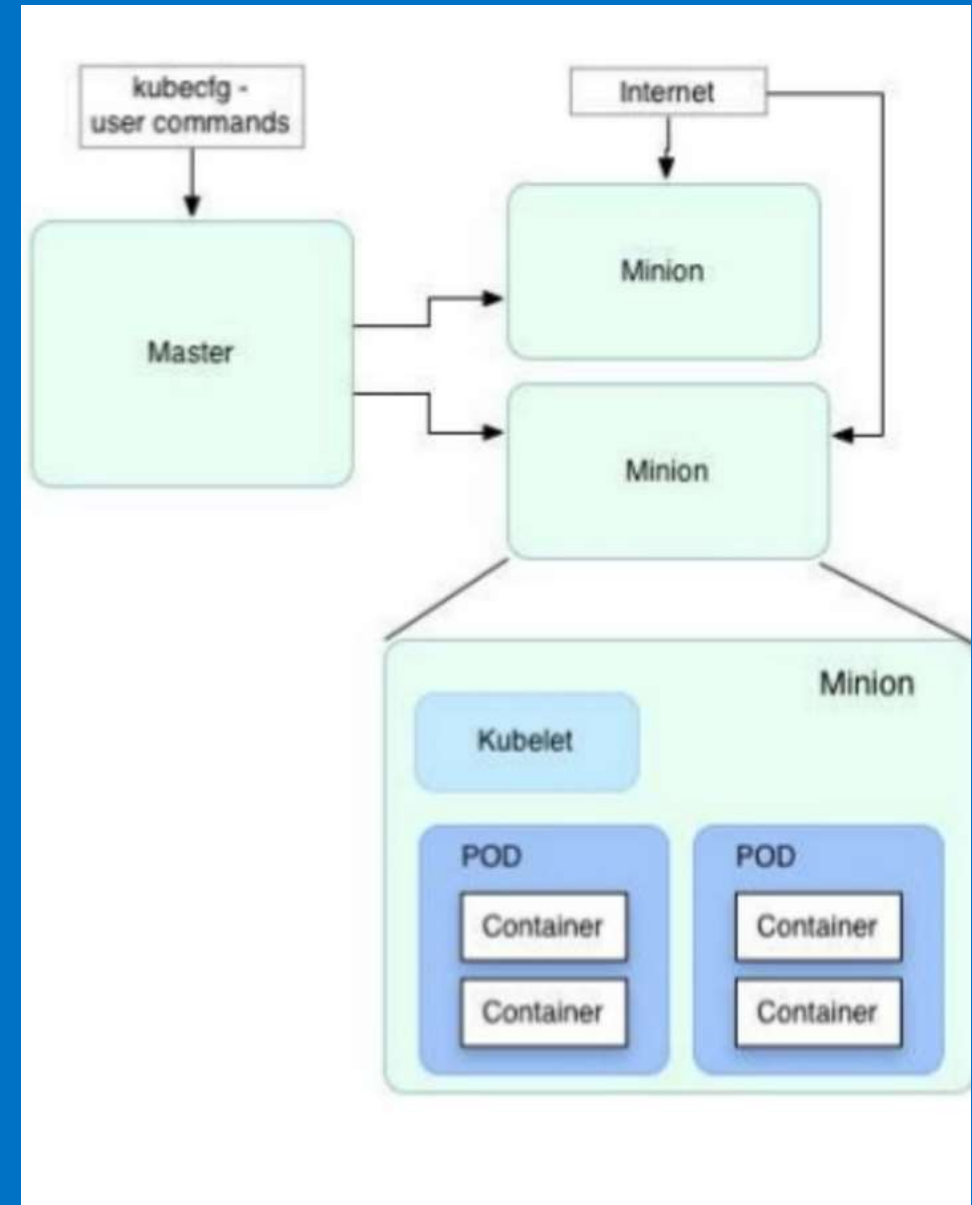
Component which runs on each agent and manages the pod and container lifecycle
There is 1:1 mapping between a host and a kubelet

Key elements of a kubelet

- Docker client
- Root Directory
- Pod Workers
- Etc client
- cAdvisor client

Functions performed by kubelet


- Run containers in pods
- Bind volumes to pods
- Bind ports to container
- Kill container
- Collect health information for container



Pod

Pod is a set of containers that can run on a host

- Pod represents a logical construct to bundle one or more applications together
 - Sidecar Model
 - Pod resources deployed to a single agent
- An application-specific "logical host"
- Share:
 - Storage resources
 - Unique network IP
- Options to control how container(s) should run
- Most commonly Docker, but other runtimes are supported



Application Health and Readiness

- Kubelets on workers use liveness and readiness probes to know when to restart a container or start directing traffic
- Liveness check: when to restart
- Readiness check: when to forward Service traffic to a pod

Labels & Selectors

- Labels
 - **Identifying** attributes on kubernetes objects
 - Key/value based

```
labels:  
  name: calcbackend-pod  
  environment: qa
```

Can be used for manual query and will be used by e.g. services

> *kubectrl get pods -l environment=production,tier=frontend*


Service

- Abstraction which defines a logical set of pods

Mortality of a single pod does **not** affect the availability of a service if the functionality of the service is provided by multiple (and not just one) pods

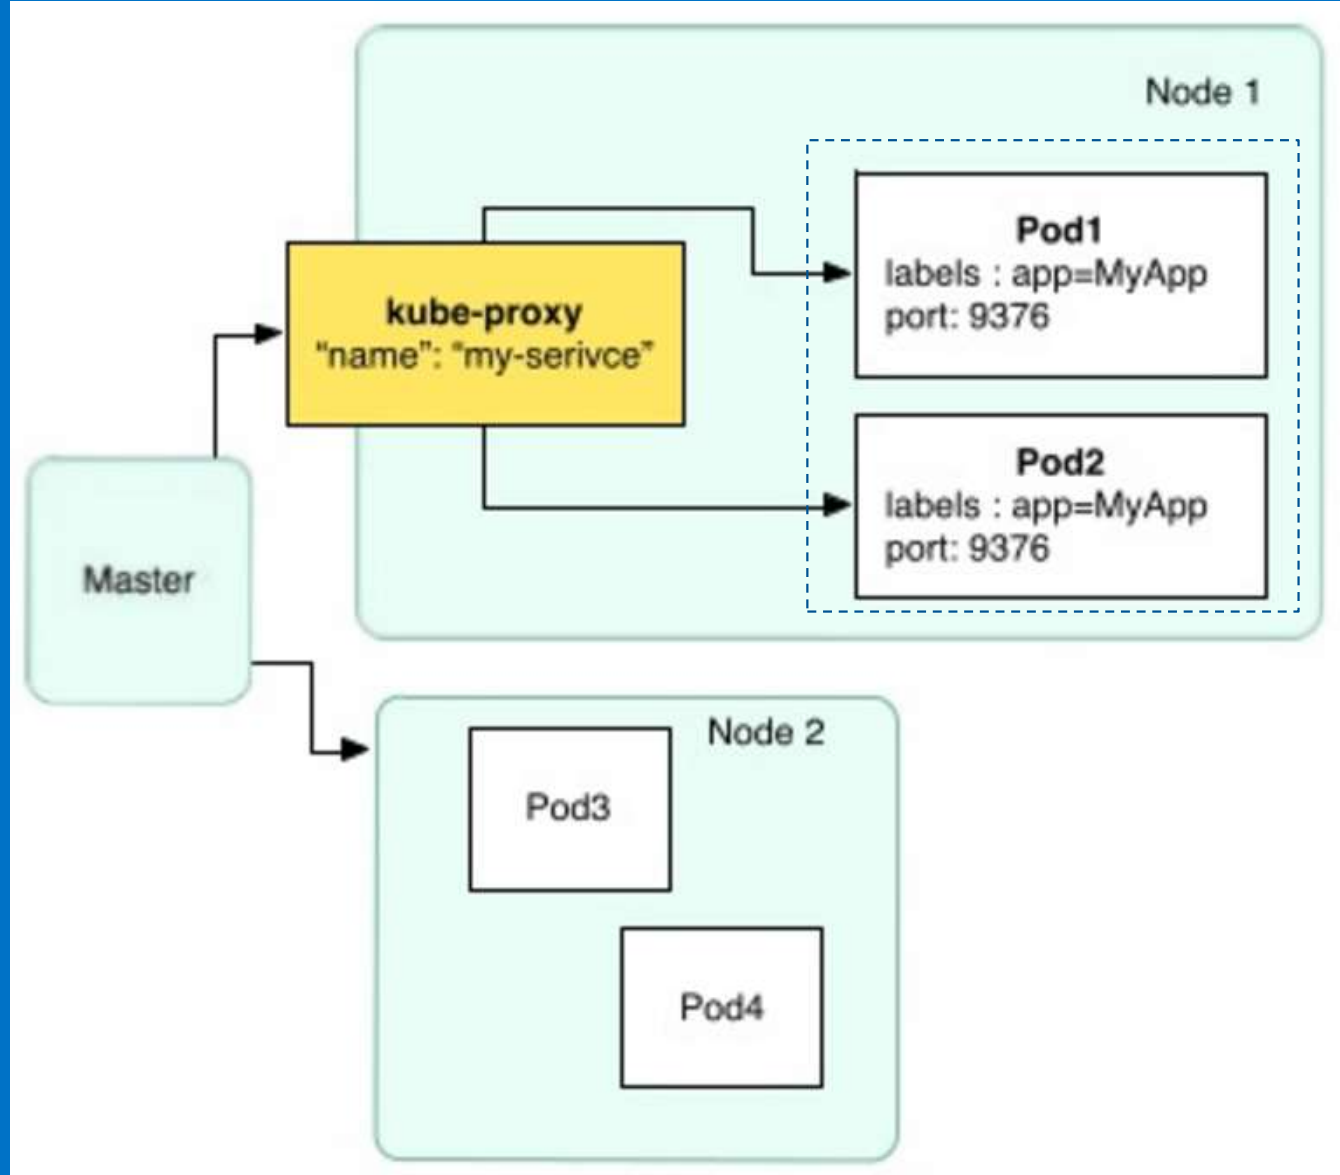
- Can be used as external interface with public IP

- Services decide via Selectors & Labels where to route traffic



Services

```
{
  "kind": "Service",
  "apiVersion": "v1",
  "metadata": {
    "name": "my-service"
  },
  "spec": {
    "selector": {
      "app": "MyApp"
    },
    "ports": [
      {
        "protocol": "TCP",
        "port": 80,
        "targetPort": 9376
      }
    ]
  }
}
```

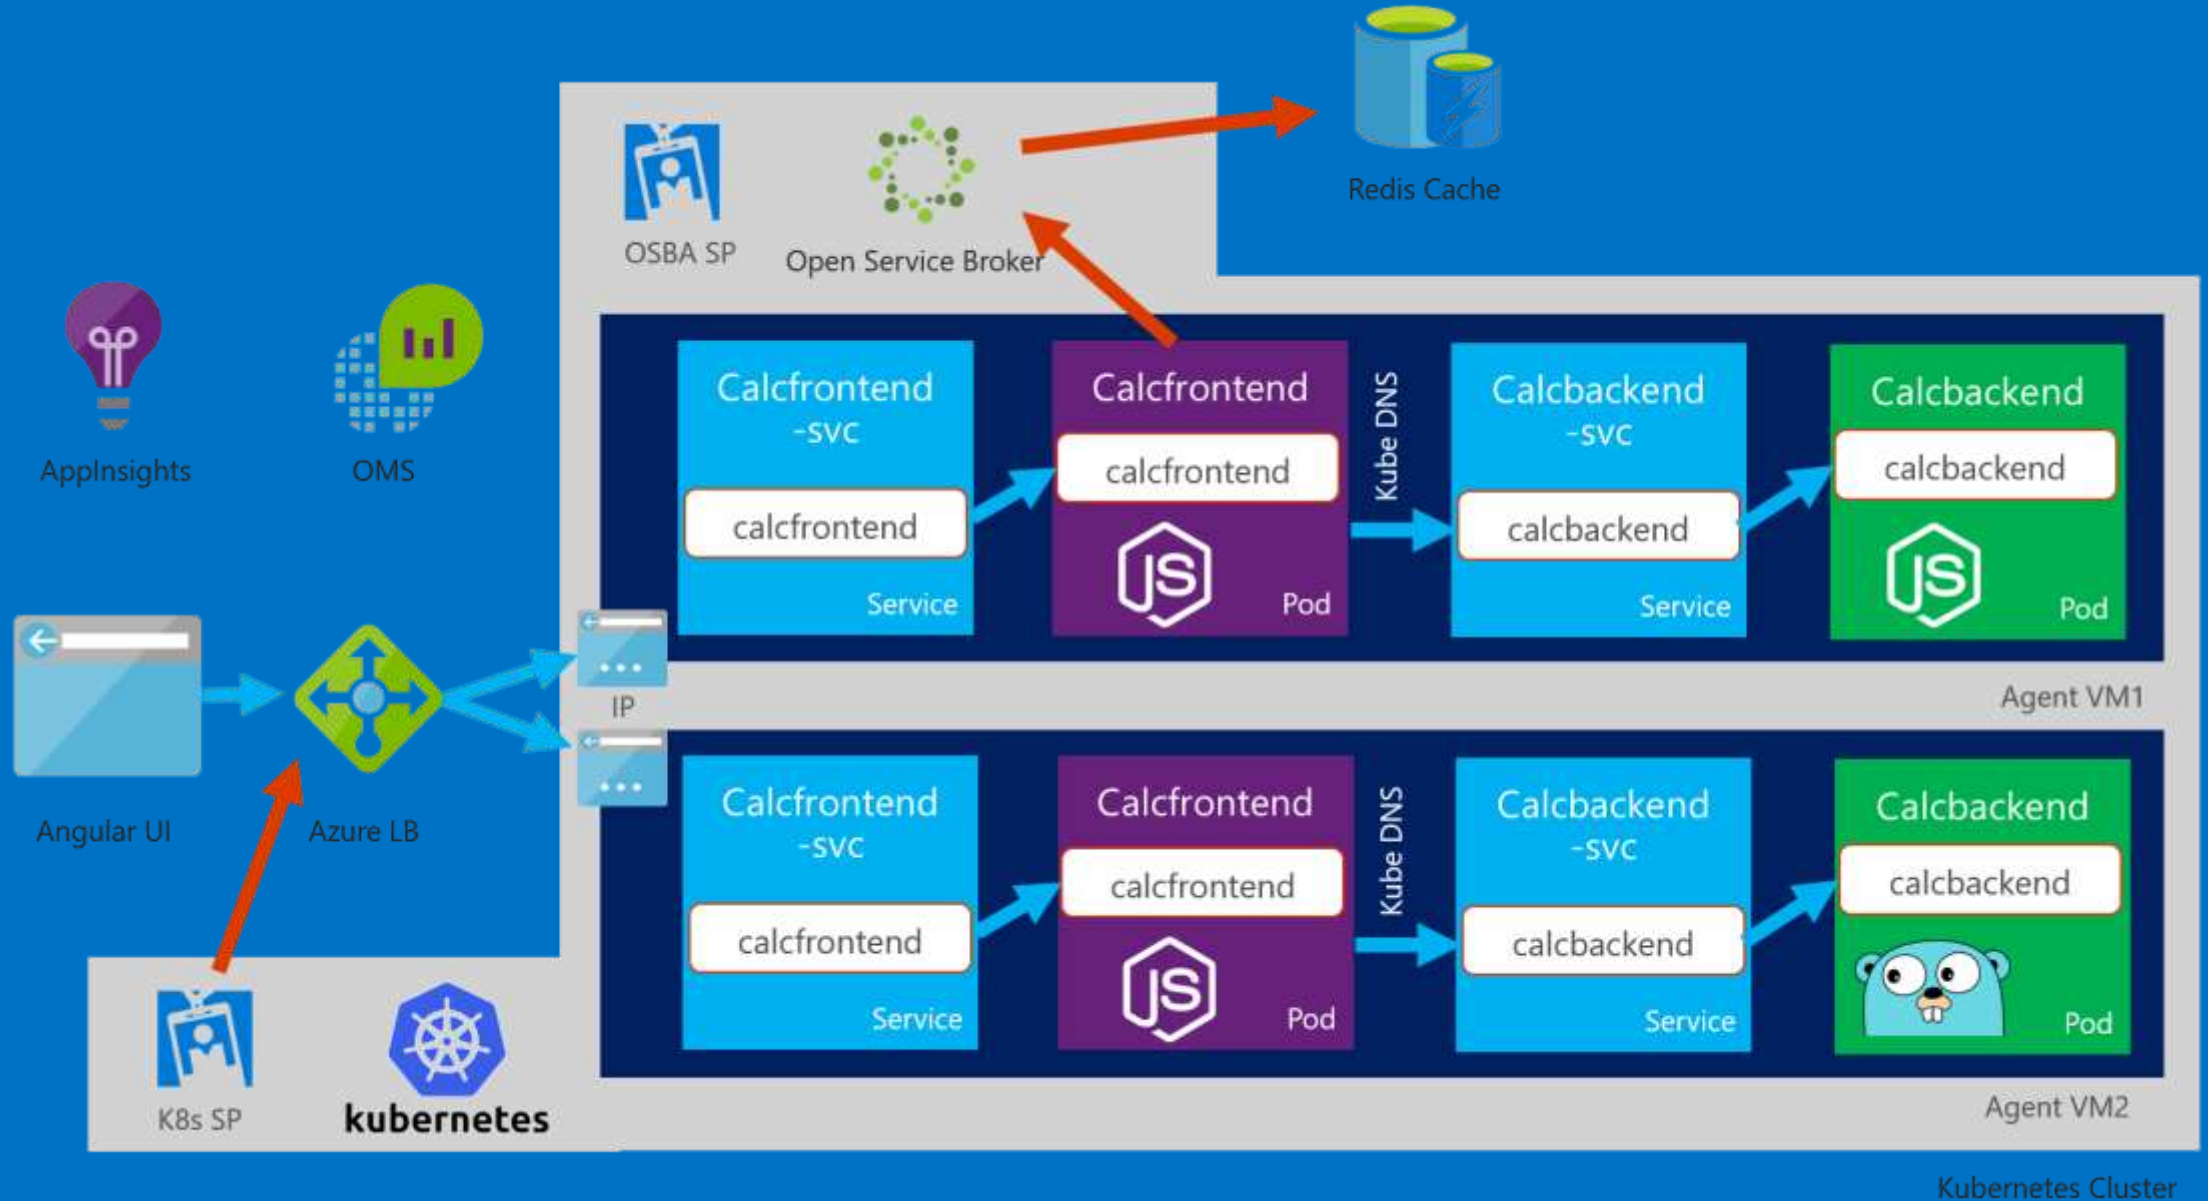


Kubernetes API


- Create, Query and modify Kubernetes objects
- E.g.
 - kubectl apply -f file.yaml
 - kubectl delete pods/mypod

```
calcbackend-pod.yml ✕
1  apiVersion: "v1"
2  kind: Pod
3  metadata:
4    name: calcbackend-pod
5    labels:
6      name: calcbackend-pod
7  spec:
8    containers:
9      - name: calcbackend-container
10        image: dmxacrmaster-microsoft.azurecr.io/calcbackend:25
11        ports:
12          - containerPort: 3000
13            name: http
14        env:
15          - name: "INSTRUMENTATIONKEY"
16            value: "abc"
17          - name: "PORT"
18            value: "3000"
```

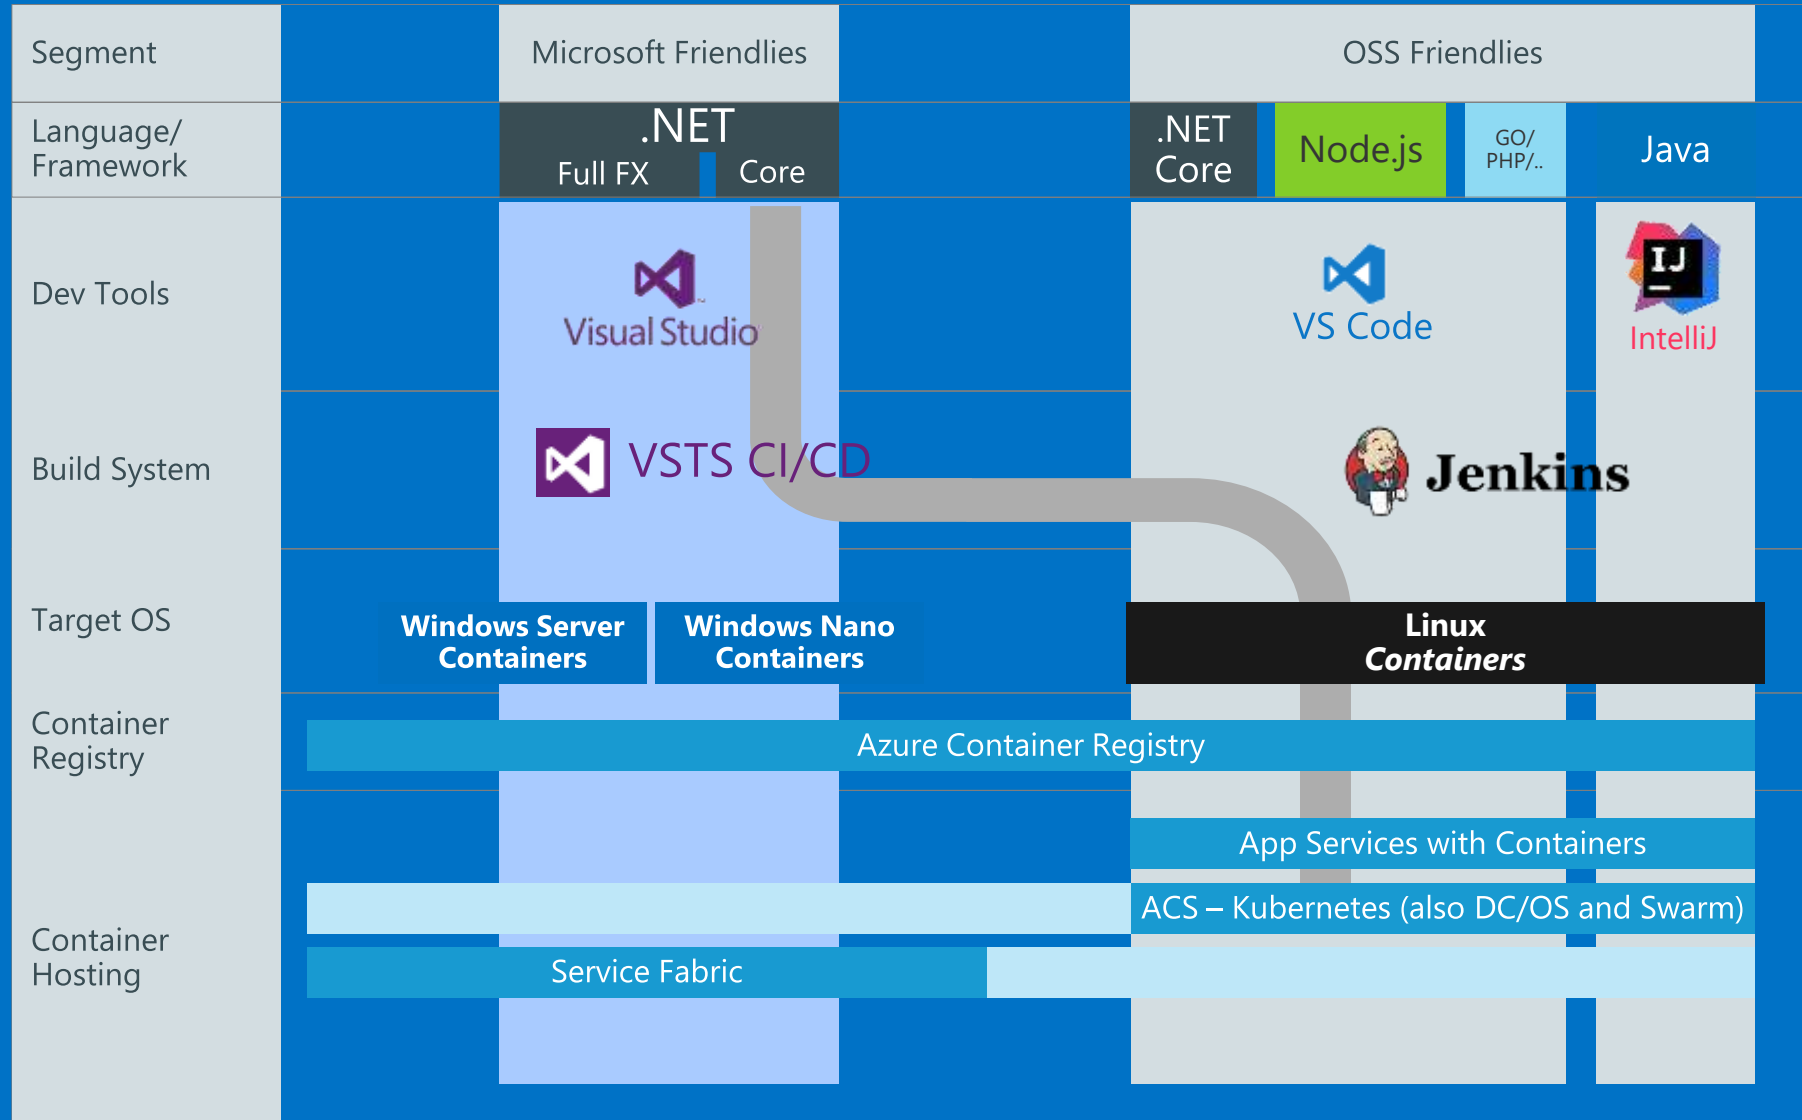
Mapping from Infrastructure & logical view



Helm, Charts & Draft




Microsoft & OSS Toolchains




Differentiation between Azure & ISV Solutions



Azure Container Instance



Container Service



Service Fabric



Web Apps



Batch




Kubernetes



Mesos DC/OS



Pivotal



OpenShift



Docker Enterprise

Thank you

